


Reading and Writing Data with Pandas



Methods to read data are all named **pd.read_*** where * is the file type. Series and DataFrames can be saved to disk using their **to_*** method.

Usage Patterns

- Use **pd.read_clipboard()** for one-off data extractions.
- Use the other **pd.read_*** methods in scripts for repeatable analyses.



Reading Text Files into a DataFrame

Colors highlight how different arguments map from the data file to a DataFrame.

```
# Historical_data.csv
Date, Cs, Rd
2005-01-03, 64.78, -
2005-01-04, 63.79, 201.4
2005-01-05, 64.46, 193.45
...
Data from Lab Z.
Recorded by Agent E
```

```
>>> read_table(
    'historical_data.csv',
    sep=',',
    header=1,
    skiprows=1,
    skipfooter=2,
    index_col=0,
    parse_dates=True,
    na_values=['-'])
```

Date	Cs	Rd


Other arguments:

- **names**: set or override column names
- **parse_dates**: accepts multiple argument types, see on the right
- **converters**: manually process each element in a column
- **comment**: character indicating commented line
- **chunksize**: read only a certain number of rows each time


Possible values of **parse_dates**:

- **[0, 2]**: Parse columns 0 and 2 as separate dates
 - **[[0, 2]]**: Group columns 0 and 2 and parse as single date
 - **{'Date': [0, 2]}**: Group columns 0 and 2, parse as single date in a column named Date.
- Dates are parsed *after* the **converters** have been applied.

Parsing Tables from the Web



```
>>> df_list = read_html(url)
```



Writing Data Structures to Disk

Writing data structures to disk:

```
> s_df.to_csv(filename)
> s_df.to_excel(filename)
```

Write multiple DataFrames to single Excel file:

```
> writer = pd.ExcelWriter(filename)
> df1.to_excel(writer, sheet_name='First')
> df2.to_excel(writer, sheet_name='Second')
> writer.save()
```

From and To a Database

Read, using SQLAlchemy. Supports multiple databases:

```
> from sqlalchemy import create_engine
> engine = create_engine(database_url)
> conn = engine.connect()
> df = pd.read_sql(query_str_or_table_name, conn)
```

Write:

```
> df.to_sql(table_name, conn)
```

Pandas Data Structures: Series and DataFrames



A Series, **s**, maps an index to values. It is:

- Like an ordered dictionary
- A Numpy array with row labels and a name

A DataFrame, **df**, maps index and column labels to values. It is:

- Like a dictionary of Series (columns) sharing the same index
- A 2D Numpy array with row and column labels

s_df applies to both Series and DataFrames.

Assume that manipulations of Pandas object return copies.

Creating Series and DataFrames

Series

```
> pd.Series(values, index=index, name=name)
> pd.Series({'idx1': val1, 'idx2': val2})
```

Where **values**, **index**, and **name** are sequences or arrays.

Series		
Values		
n1	'Cary'	0
n2	'Lynn'	1
n3	'Sam'	2
Index	Integer location	

DataFrame

	Age	Gender	Columns
'Cary'	32	M	
'Lynn'	18	F	
'Sam'	26	M	
Index	Values		

DataFrame

```
> pd.DataFrame(values, index=index, columns=col_names)
> pd.DataFrame({'col1': series1_or_seq, 'col2': series2_or_seq})
```

Where **values** is a sequence of sequences or a 2D array

Manipulating Series and DataFrames

Manipulating Columns

```
df.rename(columns={old_name: new_name})
```

Renames column

```
df.drop(name_or_names, axis='columns')
```

Drops column name

Manipulating Index

```
s_df.reindex(new_index)
```

Conform to new index

```
s_df.drop(labels_to_drop)
```

Drops index labels

```
s_df.rename(index={old_label: new_label})
```

Renames index labels

```
s_df.sort_index()
```

Sorts index labels

```
df.set_index(column_name_or_names)
```

Inserts index into columns, resets index to default integer index.

```
s_df.reset_index()
```

Manipulating Values

All row values and the index will follow:

```
df.sort_values(col_name, ascending=True)
df.sort_values(['X', 'Y'], ascending=[False, True])
```

Important Attributes and Methods

s_df.index	Array-like row labels
df.columns	Array-like column labels
s_df.values	Numpy array, data
s_df.shape	(n_rows, m_cols)
s.dtype, df.dtypes	Type of Series , of each column
len(s_df)	Number of rows

s_df.head() and s_df.tail()	First/last rows
s.unique()	Series of unique values
s_df.describe()	Summary stats
df.info()	Memory usage

Indexing and Slicing

Use these attributes on Series and DataFrames for indexing, slicing, and assignments:

s_df.loc[]	Refers only to the index labels
s_df.iloc[]	Refers only to the integer location, similar to lists or Numpy arrays

s_df.xs(key, level) Select rows with label **key** in level **level** of an object with MultiIndex.

Masking and Boolean Indexing

Create masks with, for example, comparisons

```
mask = df['X'] < 0
```

Or **isin**, for membership mask

```
mask = df['X'].isin(list_valid_values)
```

Use masks for indexing (must use **loc**)

```
df.loc[mask] = 0
```

Combine multiple masks with bitwise operators (and (&), or (|), xor (^), not (~)) and group them with parentheses:

```
mask = (df['X'] < 0) & (df['Y'] == 0)
```

Common Indexing and Slicing Patterns

rows and **cols** can be values, lists, Series or masks.

s_df.loc[rows]	Some rows (all columns in a DataFrame)
df.loc[:, cols_list]	All rows, some columns
df.loc[rows, cols]	Subset of rows and columns
s_df.loc[mask]	Boolean mask of rows (all columns)
df.loc[mask, cols]	Boolean mask of rows, some columns

Using [] on Series and DataFrames

On Series, **[]** refers to the index labels, or to a slice

s['a']	Value
s[:2]	Series, first 2 rows

On DataFrames, **[]** refers to columns labels:

df['X']	Series
df[['X', 'Y']]	DataFrame
df['new_or_old_col'] = series_or_array	

EXCEPT! with a slice or mask.

df[:2]	DataFrame, first 2 rows
df[mask]	DataFrame, rows where mask is True

NEVER CHAIN BRACKETS!

✗	> df[mask]['X'] = 1	SettingWithCopyWarning
✓	> df.loc[mask, 'X'] = 1	

ENTHOUGHT

Take your Pandas skills to the next level! Register at www.enthought.com/pandas-mastery-workshop

© 2019 Enthought, Inc., licensed under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License.
To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc-nd/4.0/>

Computation with Series and DataFrames



Pandas objects do not behave exactly like Numpy arrays. They follow three main rules (see on the right). Aligning objects on the index (or columns) before calculations might be the most important difference. There are built-in methods for most common statistical operations, such as **mean** or **sum**, and they apply across one-dimension at a time. To apply custom functions, use one of three methods to do tablewise (**pipe**), row or column-wise (**apply**) or elementwise (**applymap**) operations.

The 3 Rules of Binary Operations

Rule 1:

Operations between multiple Pandas objects implement auto-alignment based on index first.

Rule 2:


Mathematical operators (+ - * / exp, log, ...) apply element by element, on the values.

Rule 3:

Reduction operations (mean, std, skew, kurt, sum, prod, ...) are applied column by column by default.

Rule 2: Element-By-Element Mathematical Operations


Rule 1: Alignment First




Use **add**, **sub**, **mul**, **div**, to set fill value.

Rule 3: Reduction Operations

`>>> df.sum()` → Series



Operates across rows by default (**axis=0**, or **axis='rows'**).
Operate across columns with **axis=1** or **axis='columns'**.



Apply a Function to Each Value

Apply a function to each value in a Series or DataFrame

`s.apply(value_to_value)` → Series

`df.applymap(value_to_value)` → DataFrame

Apply a Function to Each Series

Apply **series_to_*** function to every column by default (across rows):

`df.apply(series_to_series)` → DataFrame

`df.apply(series_to_value)` → Series

To apply the function to every row (across columns), set **axis=1**:

`df.apply(series_to_series, axis=1)`

Apply a Function to a DataFrame

Apply a function that receives a DataFrame and returns a DataFrame, a Series, or a single value:

`df.pipe(df_to_df)` → DataFrame

`df.pipe(df_to_series)` → Series

`df.pipe(df_to_value)` → Value

What Happens with Missing Values?

Missing values are represented by **NaN** (not a number) or **NaT** (not a time).

- They propagate in operations across Pandas objects (**1 + NaN → NaN**).
- They are ignored in a "sensible" way in computations, they equal 0 in **sum**, they're ignored in **mean**, etc.
- They stay **NaN** with mathematical operations (**np.log(NaN) → NaN**).

Take your Pandas skills to the next level! Register at www.enthought.com/pandas-mastery-workshop

© 2019 Enthought, Inc., licensed under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License.
To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc-nd/4.0/>

Plotting with Pandas Series and DataFrames



Pandas uses Matplotlib to generate figures. Once a figure is generated with Pandas, all of Matplotlib's functions can be used to modify the title, labels, legend, etc. In a Jupyter notebook, all plotting calls for a given plot should be in the same cell.

Setup

Import packages:

```
> import pandas as pd
> import matplotlib.pyplot as plt
```

Execute this at IPython prompt to display figures in new windows:

```
> %matplotlib
```

Use this in Jupyter notebooks to display static images inline:


```
> %matplotlib inline
```

Use this in Jupyter notebooks to display zoomable images inline:

```
> %matplotlib notebook
```


Parts of a Figure

An Axes object is what we think of as a "plot". It has a title and two Axis objects that define data limits. Each Axis can have a label. There can be multiple Axes objects in a Figure.



Plotting with Pandas Objects

Series




With a Series, Pandas plots values against the index:

```
> ax = s.plot()
```

When plotting the results of complex manipulations with **groupby**, it's often useful to **stack/unstack** the resulting DataFrame to fit the one-line-per-column assumption (see Data Structures cheatsheet).


Dataframe



With a DataFrame, Pandas creates one line per column:

```
> ax = df.plot()
```

Labels



Use Matplotlib to override or add annotations:

```
> ax.set_xlabel('Time')
> ax.set_ylabel('Value')
> ax.set_title('Experiment A')
```

Pass labels if you want to override the column names and set the legend location:

```
> ax.legend(labels, loc='best')
```

Useful Arguments to plot




- **subplots=True**: one subplot per column, instead of one line
- **figsize**: set figure size, in inches
- **x** and **y**: plot one column against another




Red Panda
Ailurus fulgens


Kinds of Plots




`df.plot.scatter(x, y)`



`df.plot.bar()`



`df.plot.hist()`



`df.plot.box()`




Manipulating Dates and Times



Use a **Datetime** index for easy time-based indexing and slicing, as well as for powerful resampling and data alignment.

Pandas makes a distinction between timestamps, called **Datetime** objects, and time spans, called **Period** objects.

Timestamps vs Periods



Converting Objects to Time Objects

Convert different types, for example strings, lists, or arrays to Datetime with:

```
> pd.to_datetime(value)
```

Convert timestamps to time spans: set period "duration" with frequency offset (see below).

```
> date_obj.to_period(freq=freq_offset)
```

Creating Ranges of Timestamps

```
> pd.date_range(start=None, end=None,
                periods=None, freq=offset,
                tz='Europe/London')
```

Specify either a start or end date, or both. Set number of "steps" with **periods**. Set "step size" with **freq**; see "Frequency offsets" for acceptable values. Specify time zones with **tz**.

Frequency Offsets

Used by **date_range**, **period_range** and **resample**:

- | | |
|--------------------------|-----------------------|
| • B: Business day | • A: Year end |
| • D: Calendar day | • AS: Year start |
| • W: Weekly | • H: Hourly |
| • M: Month end | • T, min: Minutely |
| • MS: Month start | • S: Secondly |
| • BM: Business month end | • L, ms: Milliseconds |
| • Q: Quarter end | • U, us: Microseconds |
| | • N: Nanoseconds |

For more:

Lookup "Pandas Offset Aliases" or check out **pandas.tseries.offsets**, and **pandas.tseries.holiday** modules.

Save Yourself Some Pain: Use ISO 8601 Format

When entering dates, to be consistent and to lower the risk of error or confusion, use ISO format YYYY-MM-DD:

```
✗ >>> pd.to_datetime('12/01/2000') # 1st December
Timestamp('2000-12-01 00:00:00')
✗ >>> pd.to_datetime('13/01/2000') # 13th January!
Timestamp('2000-01-13 00:00:00')
✓ >>> pd.to_datetime('2000-01-13') # 13th January
Timestamp('2000-01-13 00:00:00')
```

Creating Ranges or Periods

```
> pd.period_range(start=None, end=None,
                  periods=None, freq=offset)
```

Resampling

```
> s_df.resample(freq_offset).mean()
```

resample returns a groupby-like object that must be aggregated with **mean**, **sum**, **std**, **apply**, etc. (See also the Split-Apply-Combine cheat sheet.)

Vectorized String Operations

Pandas implements vectorized string operations named after Python's string methods. Access them through the **str** attribute of string Series

Some String Methods

```
> s.str.lower()      > s.str.strip()
> s.str.isupper()    > s.str.normalize()
> s.str.len()        and more...
```

Index by character position:

```
> s.str[0]
```

True if regular expression pattern or string in Series:

```
> s.str.contains(str_or_pattern)
```

Splitting and Replacing

split returns a Series of lists:

```
> s.str.split()
```

Access an element of each list with **get**:

```
> s.str.split(char).str.get(1)
```

Return a DataFrame instead of a list:

```
> s.str.split(expand=True)
```

Find and replace with string or regular expressions:

```
> s.str.replace(str_or_regex, new)
> s.str.extract(regex)
> s.str.findall(regex)
```

Combining DataFrames



Tools for combining Series and DataFrames together, with SQL-type joins and concatenation. Use join if merging on indices, otherwise use merge.

Concatenating DataFrames

```
> pd.concat(df_list)
```

"Stacks" DataFrames on top of each other.

Set `ignore_index=True`, to replace index with `RangeIndex`.

Note: Faster than repeated `df.append(other_df)`.

Join on Index

```
> df.join(other)
```

Merge DataFrames on indexes. Set `on=columns` to join on index of `other` and on `columns` of `df`. `join` uses `pd.merge` under the covers.

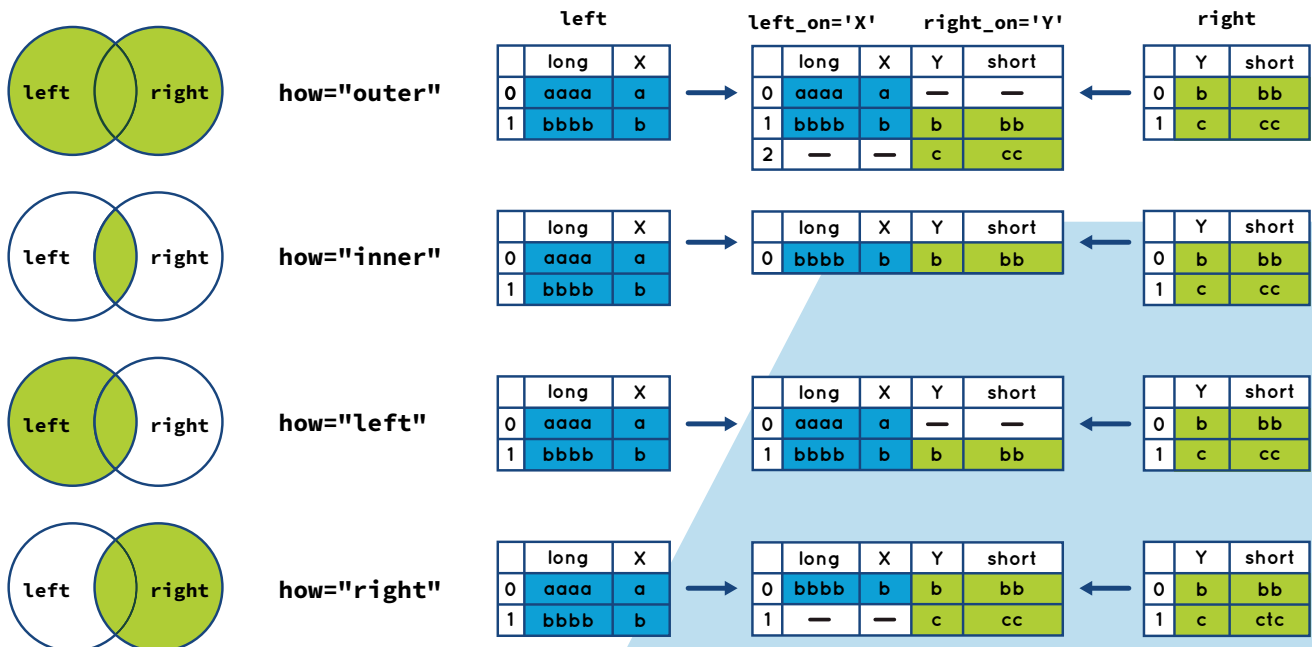
Merge on Column Values

```
> pd.merge(left, right, how='inner', on='id')
```

Ignores index, unless `on=None`. See value of `how` below.

Use `on` if merging on same column in both DataFrames, otherwise use `left_on, right_on`.

Merge Types: The `how` Keyword



Cleaning Data with Missing Values

Pandas represents missing values as `NaN` (Not a Number). It comes from Numpy and is of type `float64`. Pandas has many methods to find and replace missing values.

Replacing Missing Values

Find Missing Values

```
> s_df.isnull() or > pd.isnull(obj)
> s_df.notnull() or > pd.notnull(obj)
```

```
s_df.loc[s_df.isnull()] = 0
```

 Use mask to replace `NaN`

```
s_df.interpolate(method='linear')
```

 Interpolate using different methods

```
s_df.fillna(method='ffill')
```

 Fill forward (last valid value)

```
s_df.fillna(method='bfill')
```

 Or backward (next valid value)

```
s_df.dropna(how='any')
```

 Drop rows if any value is `NaN`

```
s_df.dropna(how='all')
```

 Drop rows if all values are `NaN`

```
s_df.dropna(how='all', axis=1)
```

 Drop across columns instead of rows

ENTHOUGHT

6

Take your Pandas skills to the next level! Register at www.enthought.com/pandas-mastery-workshop

© 2019 Enthought, Inc., licensed under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License.
To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc-nd/4.0/>


Split / Apply / Combine with DataFrames



1. *Split* the data based on some criteria.
2. *Apply* a function to each group to aggregate, transform, or filter.
3. *Combine* the results.

The apply and combine steps are typically done together in Pandas.

Split/Apply/Combine



Split: Group By

Group by a single column:


```
> g = df.groupby(col_name)
```

Grouping with list of column names creates DataFrame with MultiIndex. (see "Reshaping DataFrames and Pivot Tables" cheatsheet):

```
> g = df.groupby(list_col_names)
```

Pass a function to group based on the index:

```
> g = df.groupby(function)
```



Apply/Combine: General Tool: apply

More general than **agg**, **transform**, and **filter**. Can aggregate, transform or filter. The resulting dimensions can change, for example:

```
> g.apply(lambda x: x.describe())
```


Apply/Combine: Transformation

The shape and the index do not change.

```
> g.transform(df_to_df)
```

Example, normalization:


```
> def normalize(grp):  
    return (grp - grp.mean()) / grp.var()  
> g.transform(normalize)
```



Apply/Combine: Filtering


Returns a group only if condition is true.

```
> g.filter(lambda x: len(x)>1)
```



Other Groupby-Like Operations: Window Functions

- **resample**, **rolling**, and **ewm** (exponential weighted function) methods behave like GroupBy objects. They keep track of which row is in which "group". Results must be aggregated with **sum**, **mean**, **count**, etc. (see Aggregation).
- **resample** is often used before **rolling**, **expanding**, and **ewm** when using a DateTime index.



Split

- Groupby
- Window Functions

Apply

- Apply
- Group-specific transformations
- Aggregation
- Group-specific Filtering

Combine

Split: What's a GroupBy Object?

It keeps track of which rows are part of which group.

```
> g.groups
```

Dictionary, where keys are group names, and values are indices of rows in a given group. It is iterable:


```
> for group, sub_df in g:
```

...

Apply/Combine: Aggregation

Perform computations on each group. The shape changes; the categories in the grouping columns become the index. Can use built-in aggregation methods: **mean**, **sum**, **size**, **count**, **std**, **var**, **sem**, **describe**, **first**, **last**, **nth**, **min**, **max**, for example:

```
> g.mean()  
... or aggregate using custom function:  
> g.agg(series_to_value)  
... or aggregate with multiple functions at once:  
> g.agg([s_to_v1, s_to_v2])  
... or use different functions on different columns.  
> g.agg({'Y': s_to_v1, 'Z': s_to_v2})
```



Reshaping DataFrames and Pivot Tables



Tools for reshaping **DataFrames** from the wide to the long format and back. The long format can be tidy, which means that "each variable is a column, each observation is a row". Tidy data is easier to filter, aggregate, transform, sort, and pivot. Reshaping operations often produce multi-level indices or columns, which can be sliced and indexed.

1 Hadley Wickham (2014) "Tidy Data", <http://dx.doi.org/10.18637/jss.v059.i10>

MultiIndex: A Multi-Level Hierarchical Index

Often created as a result of:

```
> df.groupby(list_of_columns)
> df.set_index(list_of_columns)
```

Contiguous labels are *displayed* together but apply to each row. The concept is similar to multi-level columns.

A **MultiIndex** allows indexing and slicing one or multiple levels at once. Using the *Long* example from the right:

```
long.loc[1900]    All 1900 rows
long.loc[(1900, 'March')]    value 2
long.xs('March', level='Month')    All March rows
```

Simpler than using boolean indexing, for example:

```
> long[long.Month == 'March']
```

Long to Wide Format and Back with **stack()** and **unstack()**

Pivot **column** level to index, i.e. "stacking the columns" (wide to long):

```
> df.stack()
```

Pivot **index** level to columns, "unstack the columns" (long to wide):

```
> df.unstack()
```

If multiple indices or column levels, use level number or name to **stack/unstack**:

```
> df.unstack(1) or > df.unstack('Month')
```

A common use case for unstacking, plotting group data vs index after groupby:

```
> (df.groupby(['A', 'B'])['relevant'].mean()
   .unstack().plot())
```

Wide				Long			
Year	Jan.	Feb.	Mar.	Year	Month	Value	
1900	1	7	2	1900	Jan.	1	
2000	4	3	9	2000	Feb.	7	
					Mar.	2	
					Jan.	4	
					Feb.	3	
					Mar.	9	

Pivot Tables

```
> pd.pivot_table(df,
  index=cols, (keys to group by for index)
  columns=cols2, (keys to group by for columns)
  values=cols3, (columns to aggregate)
  aggfunc='mean') (what to do with repeated values)
```

Omitting index, columns, or values will use all remaining columns of df. You can "pivot" a table manually using **groupby**, **stack** and **unstack**.

Index			Columns		
0	Recently updated	Number of stations	Continent code	AN	EU
1	FALSE	1	Recently updated		
2	FALSE	1	FALSE	1	3
3	FALSE	1	TRUE	2	1
4	TRUE	1			
5	FALSE	1			
6	TRUE	1			
7	TRUE	1			

```
pd.pivot_table(df,
  index="Recently updated",
  columns="continent code",
  values="Number of Stations",
  aggfunc=np.sum)
```

From Wide to Long with **melt**

Specify which columns are identifiers (**id_vars**, values will be repeated for each row) and which are "measured variables" (**value_vars**, will become values in *variable* column. All remaining columns by default).

```
pd.melt(df, id_vars=id_cols, value_vars=value_columns)
```

```
pd.melt(team, id_vars=['Color'],
  value_vars=['A', 'B', 'C'],
  var_name='Team', value_name='Score')
```

Team				Melt			
Color	A	B	C	Color	Team	Score	
Red	1	3	4	Red	A	1	
Blue	2	-	6	Blue	A	2	
				Red	B	3	
				Blue	B	-	
				Red	C	4	
				Blue	C	6	

df.pivot() vs pd.pivot_table

df.pivot() Does not deal with repeated values in index. It's a declarative form of **stack** and **unstack**.

pd.pivot_table() Use if you have repeated values in index (specify **aggfunc** argument).



Red Panda
Ailurus fulgens